Does My Toaster Get Rights?

While playing the game Fallout 4 I ran into two groups trying to accomplish conflicting

tasks. The Brotherhood of Steel was fighting to exterminate the underground organization known

as The Institute and all of their Synths. The Railroad on the other hand believed that Synths had

the same rights as humans and should be saved from the Institute and be allowed to live their

own lives. Synths are androids that the Institute created to do menial labor and be soldiers in

their armies. Synths are categorized by one of three generations. Generation one Synths have

bare minimum AI and only have a metallic skeleton for a body, they are used only for basic grunt

work and as expendable soldiers. Generation two Synths have a more advanced AI allowing

them to think more like humans and have some semblance of free will and emotions, their bodies

look like mannequins and are covered with flesh colored rubber. Generation three Synths look

exactly like humans and are indistinguishable by sight, their AIs are capable of complex

thoughts, emotions and free will to the extent that humans are. The Brotherhood of Steel believes

that all the Synths should be destroyed irregardless of whether or not they wish to be free and

live regular human lives, their philosophy is that since they are robots they do not deserve the

same rights as humans. The Railroad believes that having rights is not specific to humans and

therefore Synths are capable of deserving rights. So, who is right? Should we give rights to

artificial intelligence or are they still tools that can be used by humanity no matter their qualities?

The Railroad's way of thinking about AI rights is the most ethical way to think no matter what

system of ethics you apply; the real question is not a matter of if AI deserve rights, it is a matter

of where we draw the line between being a tool and deserving of rights.

Before we begin any type of discussion on what rights AIs deserve or do not deserve we have to first establish that AIs are capable of even having rights in the first place. Some people will take the Brotherhood of Steel's side and say that because AI are not human they do not deserve any rights. If this view is true, that means the only creatures that may be granted rights are humans. This approach to who has rights and who does not contains many inconsistencies and if actually applied, leads to many undesirable conclusions. Right from the get-go we arrive at the issue that the definition for being human has never been clearly defined and is actually an extremely large part of the abortion argument, so for the sake of this paper and simplicity "human" will be defined as something containing human DNA. Taking this definition, we must say that whenever we get a haircut we leave behind humans on the floor of the salon and that they deserve the same rights as the body they came from because they contain human DNA and they are separate from everyone else. Since no one is willing to acknowledge that their discarded hair deserves rights this definition is clearly invalid. Even if we revise the definition to say that only living creatures with human DNA can be considered human we still end up with having cancer cells being humans and deserving of rights. Cancer cells are cells in a person's body that have a slightly different DNA than their host which results in them multiplying according to their own needs and not the host's and many times killing their host. Currently we do not think of cancer cells as humans and try to kill the cells if it would result in the death of the host, but if we assume any living thing containing human DNA is a human and deserves rights we may not kill the cancer cells in order to save the host because their lives are both equally important.

In addition to these issues, stating that only humans are deserving of rights leads to the conclusion that all non-humans have no rights. It follows then that animals have no rights and

that we may treat them with as much cruelty as we desire without violating any ethical rules. This is obviously not the case as the good majority of people believe that animals deserve some form of rights even though they are non-human creatures. Most ethical theories also support the idea that non-human creatures deserve rights to a certain extent, therefore all the theories that support animal rights also work for supporting the rights of AI. Utilitarianism bases its ethical rules off the creation of happiness and pain, for utilitarians any action that causes more pain than happiness is considered unethical. Utilitarian philosophers such as Peter Singer argue that harming animals, including eating them, is unethical because it causes suffering and all suffering is equal. If an AI is able to determine for itself what actions it wishes to happen to it and what it wishes to avoid then we must say any action that an AI wishes to avoid that is acted against it causes the AI suffering, and is therefore unethical under utilitarianism. After all, what is suffering other than having something happen to you that you do would rather have avoided? Even Carl Cohen, a philosopher who argues against animal rights, only does so because he claims that they are not deserving of rights because they cannot make moral claims and they can not respond to claims made against them. This is not the case with AI, gen 3 Synths are capable of wanting certain things and being upset with humans who treat them in ways they believe they do not deserve to be treated. In this way these Synths are able to make moral claims against humans and they are also capable of understanding moral claims made against them and they can choose to follow them or ignore them. Under Cohen's philosophy AIs are not only capable of rights but should be given the exact same rights as humans because they are capable of making moral claims and having moral claims made against them just like every human.

The final argument against AI deserving any rights is that they are unnatural, they are not created from cells but rather machinery. This argument is laughably inadequate because it draws the line for the capacity for rights not at a meaningful place like self awareness but at what the entity is composed of. Anyone supporting this view would have to say that if one day we managed to put an AI into a body made out of living cells that AI would have human rights while the rest did not, and this is not a valid way to go about assigning rights.

Now the question switches from whether or not AI are capable of rights to how many rights we give to AI and how we should treat them. Assuming that we all agree that AI are capable of having rights, most people would agree that AI like gen 3 Synths deserve the same rights as humans. This is because they think and act just as humans do so they are identical to us in every way except for their composition. The justification for giving gen 3 Synths the same rights as humans is pretty straightforward so we will instead focus on what rights gen 2 and gen 1 Synths are given. What rights we give to who largely depends on which ethical theory we choose to subscribe to therefore there will be multiple analyses.

From the utilitarian point of view we must do our best to reduce suffering where and when we can, and the most ethical choice in a situation is one that has the greatest net happiness. This way of determining what is right and wrong raises issues but that is not the point of this argument, what matters is that the suffering of AI and humans can not be said to be significantly different enough to say humans deserve more of a say in the utilitarian calculation. What is suffering other than having something happen to you that you do not wish to happen? This is a very simplistic view of suffering but not one that is necessarily incorrect due to the fact that we can not say someone is suffering from an event if they wanted it to occur and enjoy the outcome.

This is of course assuming that the someone in question is mentally sane. If we take the fact human suffering and the suffering of AI are considered equal and our definition of suffering, we must conclude that any AI with the capability to desire to avoid certain situations. This is not to say that simply acknowledging that circumstances should be avoided is the same as disliking said circumstances. It is relatively simple to code an AI that can take in a bit of data and use some machine learning magic to figure out what people generally view as enjoyable and unenjoyable experiences but this does not mean that the AI wants either of these simply that it can answer a question based on data. This sort of answer from the AI is no different than the facial recognition software Facebook uses to suggest tags, we don't say that Facebook really knows us and our friends so we wouldn't say that an AI this simple wants either situation. However if an AI is capable of understanding the consequences of the situation it is presented with and it wishes to avoid it we should consider the AI as suffering if the situation occurs. In Fallout 4 gen 1 Synths are only capable of understanding their orders and calculating how to complete their tasks so utilitarians would not consider anything done to them as unethical. The other two generations however are capable of understanding their existence and wanting to preserve themselves so certain actions towards these groups would cause an increase in suffering.

Virtue ethics focuses more on how we should behave and what type of person we should strive to become rather than trying to categorize every situation into ethical and unethical. It is because of virtue ethics' introspective nature that it does nothing to say what rights should apply to AI. Instead virtue ethics tells every individual to follow a set of virtues to become the best they can be and in doing so they would have to be kind to AI. Since some AI are capable of wanting and not wanting certain situations and understanding the world around them and adapting, under

virtue ethics we must consider all the virtues when interacting with them. Subscribing to virtue ethics has another effect which many would consider decently uplifting. Since AIs are capable of understanding the world around them they would also be able to be taught these virtues and act accordingly.

In a piece against abortion, Mary Anne Warren stated that only persons have moral rights and because fetuses can not be considered persons they do not have the same rights as persons. Warren goes on to define a person as someone who: is sentiment and self aware, emotional and rational, capable of communication, and capable of moral agency. It is sad to think that under this definition an AI has more rights than a human fetus, but not if you think about the context they are taken in. An AI does not grow like a human mind, they are programmed with the ability to learn and improve they are not forced to start at the same level as a human baby each time they are created. The AIs we are discussing in this paper have the mental capacities of adult humans, though depending on the complexity one could argue that they are not quite there. Whether or not they are completely passable for adult humans is not important because they are capable of all the requirements needed to fulfill every definition of moral responsibility, capacity to feel pain, and capacity to make moral claims and have moral claims made against them. Even if someone were to argue that these could be taken away from an AI because all they are is code and we can alter that, the exact same can be done for humans. All we are are brains piloting a meat suit, and if we were to remove a portion of someone's brain they would be able to continue living but would lose the mental abilities to complete many of the tasks required by different philosophers to deserve rights. This does not mean that if someone has been lobotomized, or disabled for that matter, that they suddenly have no rights and we may do whatever we wish to

them. Due to the fact that no reasonable person would say once you become mentally impaired to the point of not fitting the definition of a person you lose your rights, we can comfortably say that rights are not something that can be lost once gained.

In conclusion if we ever create artificial intelligence that is capable of understanding the consequences of its actions and the actions of others we must grant it the same rights that we grant humans. This is because this AI would be able to understand the world close enough to how we do that it can make moral claims and we can make moral claims against it, it can also choose not to act immorally towards us based on those claims.